

Readers/Writers problem

- There is a number of threads.
 - In order to ensure the integrity of the shared data being both read from and written to we need to allow:
 - only one writer access to the data at a time
 - if a writer is active there must be no active readers
 - if no writer is active there can be multiple readers
- We also need to make sure that no process misses out entirely.
- Three types of solutions:
 1. writer preferred - waiting writers go before waiting readers
 2. reader preferred – waiting readers go before waiting writers
 3. neither preferred - try to treat readers and writers fairly (a simple queue is not good enough we want parallel readers whenever possible)
- Both 1 and 2 can lead to indefinite postponement.
- See `pthread_rwlock` (`rdlock` and `wrlock`) - which solution do they provide?

Getting the program correct

```
exclusive_access = Semaphore.new(1)
number_deposited = Semaphore.new(0)
```

```
shared_buffer = 0
```

```
producer = Thread.new do
  while true
    next_result = whatever
    exclusive_access.wait()
    shared_buffer = next_result
    number_deposited.signal()
    exclusive_access.signal()
  end
end
```

```
consumer = Thread.new do
  while true
    exclusive_access.wait()
    number_deposited.wait()
    next_result = shared_buffer
    exclusive_access.signal()
    puts next_result
  end
end
```

Programming using low level constructs like semaphores is prone to mistakes. What is wrong with this semaphore solution to the producer/consumer problem? 2 different problems.

Bad programmers

- **Another popular problem is forgetting to unlock or signal.**


We want an automatic (more or less) way of helping programmers lock and unlock.

- Java, Ruby and other languages try to avoid or minimise problems by implementing a form of monitor.

Monitors

- Brinch Hansen (1973) Hoare (1974)
- You can think of a monitor as an object which only allows one thread to be executing inside it.
It has:
 - the shared resource - it can only be accessed by the monitor
 - publicly accessible procedures - they do the work
 - a queue to get in
 - a scheduler - which thread gets access next
 - local state - not visible externally except via access procedures
 - initialization code
 - condition variables

Monitors (cont.)



Example monitor

```
monitor Account
```

```
    money = 0.00 # the shared resource
```

```
    def deposit(amount)
        money = money + amount
    end
```

```
    def withdraw(amount)
        if (amount < money)
            money = money - amount
            return true
        else
            return false
        end
    end
```

```
end
```

```
def balance
    return money
end
```

```
end
```

Here is an example in some Pseudo-code language which includes monitors:

Condition variables

- But sometimes our threads have to wait for some condition.
- A condition variable is a queue which can hold threads. We have wait and signal operations on condition variables.
- `conditionVariable.wait` puts the current thread to sleep on the corresponding queue
- `conditionVariable.signal` wakes up one thread from the queue (if there are any waiting)
- No internal state is kept of how many signals and waits there have been.
- Simpler than the similar instructions on semaphores.
- A signal with nothing waiting does nothing.
- A wait always puts a thread to sleep.

e.g. condition variables

```
monitor SimpleBuffer
```

```
    def initialize
      buffer_free = true
      buffer = 0
      empty = new_condition_var
      full = new_condition_var
    end
```

```
    def insert(value)
      while !buffer_free
        empty.wait
      end
      buffer = value
      buffer_free = false
      full.signal
    end
```

```
    def retrieve
      while buffer_free
        full.wait
      end
      data = buffer
      buffer_free = true
      empty.signal
      return data
    end
```


```
end
```


But which thread runs?

- But doesn't signal mean we have two threads running in the monitor?
- Two choices:
 - stop the thread which called signal
 - don't start the new one until the current thread leaves the monitor
- Usually we use the second answer but:
 - the thread may signal on other condition variables as well and we have to make scheduling decisions
 - it may also change the conditions again and the next thread shouldn't really run

Java monitors

- Java has a single lock variable per object (it also has one per class).
- Each object also has a *wait set* associated with it (carefully not called a queue).
- Synchronized methods and blocks must check this variable before allowing entry.



Java monitors (cont.)

- There is a count associated with each lock variable.
- The count goes up every time a thread which owns the lock on that object calls a synchronized method or block on that object.
- And it goes down when it leaves the method or block.
- When the count gets to zero the thread exits the monitor and the lock is released.

...

```
synchronized (anObject) {  
    do things to the object;  
}
```

Java monitors are different

- `signal` is called `notify()`.
- It doesn't provide condition variables in the language (but (1.5 and later) provides them as classes).
- `wait()` and `notify()` have a single set for the whole object, i.e. one condition variable.
- The object can have unsynchronized methods which are not private.
- Also fields which are not private. Not a good idea.
- after a `notify()` running threads run till they leave the synchronized area
- programmers are told to use a `while` loop with the conditional `wait`

Before next time

- Read from the textbook
 - 7.1.3 The Dining-Philosophers Problem
 - 8.3 Deadlock Characterization
 - 8.4 Methods for Handling Deadlock
 - 8.5 Deadlock Prevention
 - 8.6 Deadlock Avoidance