

Virtual memory & MMU I

CompSys304 – Computer Architecture

A/Prof Oliver Sinn

Department of Electrical, Computer, and Software Engineering

1 Principle

2 Address translation

Motivation or problems solved

- ① Safe and efficient sharing of memory among multiple programs
 - Secure sharing
- ② Burden of programming for small, limited memories
 - Programs can exceed size of physical memory

Virtual memory

Motivation or problems solved

- 1 Safe and efficient sharing of memory among multiple programs
 - Secure sharing
- 2 Burden of programming for small, limited memories
 - Programs can exceed size of physical memory

Key elements

- Virtual address – address used by CPU
- Physical address – address in main memory
- Address translation – *mapping* virtual to physical address

Virtual memory

Motivation or problems solved

- 1 Safe and efficient sharing of memory among multiple programs
 - Secure sharing
- 2 Burden of programming for small, limited memories
 - Programs can exceed size of physical memory

Key elements

- Virtual address – address used by CPU
- Physical address – address in main memory
- Address translation – *mapping* virtual to physical address

OK, but why here, why not in OS course?

- Address translation done in CPU hardware (Memory Management Unit – MMU) and software
- Impact on CPU performance


- Each program gets own address space (i.e. area)
 - Virtual memory exists in abundance
 - No other program is allowed to access this address space
- ⇒ Protection of program's code and data from other programs

- Each program gets own address space (i.e. area)
 - Virtual memory exists in abundance
 - No other program is allowed to access this address space
 - ⇒ Protection of program's code and data from other programs
- Main memory becomes cache of secondary memory (disks)
 - Indeed very similar to cache concept

- Each program gets own address space (i.e. area)
 - Virtual memory exists in abundance
 - No other program is allowed to access this address space
 - ⇒ Protection of program's code and data from other programs
- Main memory becomes cache of secondary memory (disks)
 - Indeed very similar to cache concept
- Virtual address has usually more bits than physical address
 - Idea of concept!
 - But less is also possible (32-bit CPU with more than 4GB RAM)

Principle of address translation

- Virtual address can be mapped to any physical address or disk location
- Address translation is process of doing this



- Enables relocation
 - Load program anywhere in main memory

- Enables **relocation**
 - Load program anywhere in main memory
- Allows physical addresses not to be continuous
 - Only virtual addresses need to be

- Enables **relocation**
 - Load program anywhere in main memory
- Allows physical addresses not to be continuous
 - Only virtual addresses need to be

Possible:

- 2 different virtual addresses can map to same physical address
 - To share code between programs

1 Principle

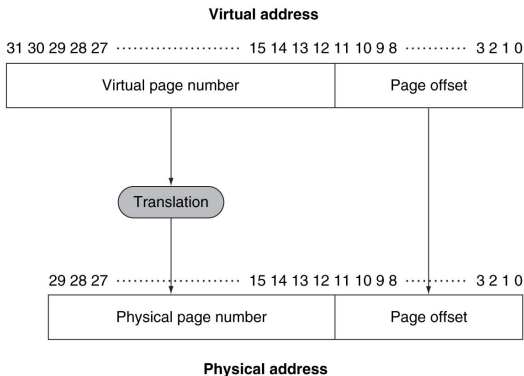
2 Address translation

Address translation notation

- **Page** – smallest block handled by virtual memory
 - Corresponds to cache line
- **Page offset f** – 2^f is size of page (in bytes)
 - On x86 CPUs (all?) $f = 12 \rightarrow$ page size 4KByte
 - Support for 2/4MB pages on x86
- **Page number** – number referencing page, index of page
- **Page hit** – page is in main memory
 - Translation handled by MMU
- **Page fault** – page is not in main memory
 - Virtual memory miss
 - Corresponds to cache miss
- **Page fault handler** – function to handle a page miss
 - SW: part of OS kernel

Address translation

- Virtual address 32 bits, physical address 30 bits, page offset 12 bits



Page fault

- Page fault is very expensive
 - Load 4KB from disk
 - Latency of random disk access very high in comparison to RAM
 - much better these days with SSDs

Page fault

- Page fault is very expensive
 - Load 4KB from disk
 - Latency of random disk access very high in comparison to RAM
 - much better these days with SSDs
- High latency
 - ⇒ Pages should not be small

- Page fault is very expensive
 - Load 4KB from disk
 - Latency of random disk access very high in comparison to RAM
 - much better these days with SSDs
- High latency
 - ⇒ Pages should not be small
- Write back (copy back)
 - Write through takes long
 - Writing single word/cache line to disk is inefficient (latency/bandwidth is large)
 - Dirty bit – set when written to any word of page

- Page fault is very expensive
 - Load 4KB from disk
 - Latency of random disk access very high in comparison to RAM
 - much better these days with SSDs
- High latency
 - ⇒ Pages should not be small
- Write back (copy back)
 - Write through takes long
 - Writing single word/cache line to disk is inefficient (latency/bandwidth is large)
 - Dirty bit – set when written to any word of page
- Page faults can be handled by software
 - Disk access slow anyway, hardware not necessary
 - Advantage, clever placement algorithms can reduce page fault rate

Placement and retrieval

- Reduce page fault rate
 - ⇒ Fully associative placement of pages in memory
 - Virtual page can go into any physical page
- But, locating a page must be fast
- Use indexed table

Placement and retrieval

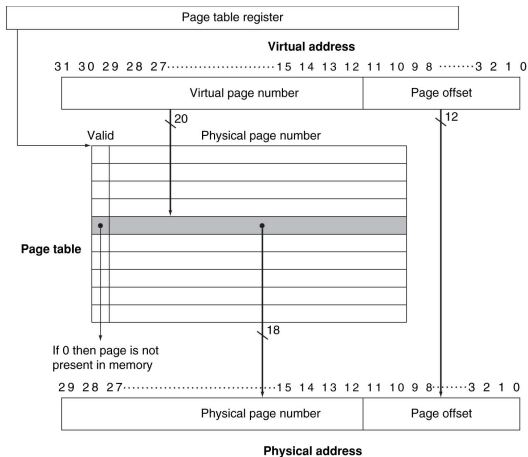
- Reduce page fault rate
 - ⇒ Fully associative placement of pages in memory
 - Virtual page can go into any physical page
- But, locating a page must be fast
- Use indexed table

Page table

- In main memory
- Every program (process) has own page table
- Virtual-physical pairs
 - also entries for not in memory (valid bit is false)
- Indexed with page number of virtual address
 - No searching necessary
- Page table register
 - MMU register: points to start of page table

Page table

- Table indexed by page number
 - $2^{20} = 1\text{M}$ entries, 4 bytes each $\Rightarrow 4\text{MB}$



Memory requirement of page tables

- On 32-bit system, page table would require 4 MB
 - Remember, one table per program (process)
 - Would consume too much memory
 - Even worse on systems with larger virtual address

Memory requirement of page tables

- On 32-bit system, page table would require 4 MB
 - Remember, one table per program (process)
 - Would consume too much memory
 - Even worse on systems with larger virtual address
- Solutions
 - E.g. Multi-level page table
 - *or* Only keep used part of table
 - Only possible to allocate consecutive parts of virtual memory
 - *or* ...

Memory requirement of page tables

- On 32-bit system, page table would require 4 MB
 - Remember, one table per program (process)
 - Would consume too much memory
 - Even worse on systems with larger virtual address
 - Solutions
 - E.g. Multi-level page table
 - *or* Only keep used part of table
 - Only possible to allocate consecutive parts of virtual memory
 - *or* ...
 - Not discussed here
- ⇒ studied in Operating System design

Replacement policy

- When physical memory exhausted, evict which page?

Replacement policy

- When physical memory exhausted, evict which page?
- LRU – Least Recently Used
 - Again, as with caches

Replacement policy

- When physical memory exhausted, evict which page?
- LRU – Least Recently Used
 - Again, as with caches
- Table contains many pages
 - Exact LRU would be too costly

Replacement policy

- When physical memory exhausted, evict which page?
- LRU – Least Recently Used
 - Again, as with caches
- Table contains many pages
 - Exact LRU would be too costly

Approximation

- Use **reference bit (ref bit)**
 - Set when page is accessed
 - Periodically reset
 - Software takes snapshots before reset
 - Can build set of recently used pages
- Evict page to swap space whose ref bit has not been set for some time