

# Virtual memory & MMU II

## CompSys304 – Computer Architecture

A/Prof Oliver Sinn

Department of Electrical, Computer, and Software Engineering

## 1 TLB

## 2 OS and Protection

# Performance problem?

Isn't there a performance problem???

# Performance problem?

## Isn't there a performance problem???

- Translating every address (by looking into a table in memory) before it is accessed
  - I.e read/write would require two memory accesses, plus translation
    - 1st page table, 2nd data at address
  - More potential for cache misses

# Performance problem?

## Isn't there a performance problem???

- Translating every address (by looking into a table in memory) before it is accessed
  - I.e read/write would require two memory accesses, plus translation
    - 1st page table, 2nd data at address
  - More potential for cache misses

Solution:

### TLB – translation lookaside buffer

- Page table cache (yes, caches again 😊)
- In MMU of every modern CPU

# Performance problem?

## Isn't there a performance problem???

- Translating every address (by looking into a table in memory) before it is accessed
  - I.e read/write would require two memory accesses, plus translation
    - 1st page table, 2nd data at address
  - More potential for cache misses

Solution:

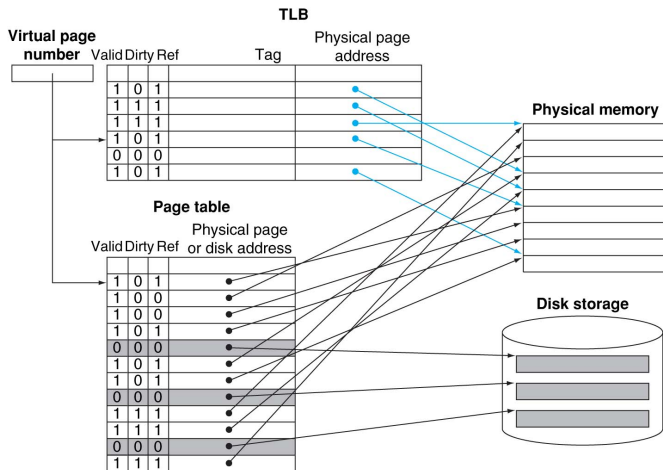
### TLB – translation lookaside buffer

- Page table cache (yes, caches again 😊)
- In MMU of every modern CPU

### When is address translation done?

- Before going to L1 cache
- ⇒ Caches use physical addresses (usually)

# TLB principle



## Metrics – typical

- Size: typical 16-512 entries
- Block size: 1-2 page table entries (4-8 bytes each)
- Hit time: 0.5-1 cycle
- Miss penalty: 10-100 cycles (page hit)
- Miss rate: 0.01%-1%
- Levels: 1-2
- Split: data and instruction split on first level or both



# TLB metrics

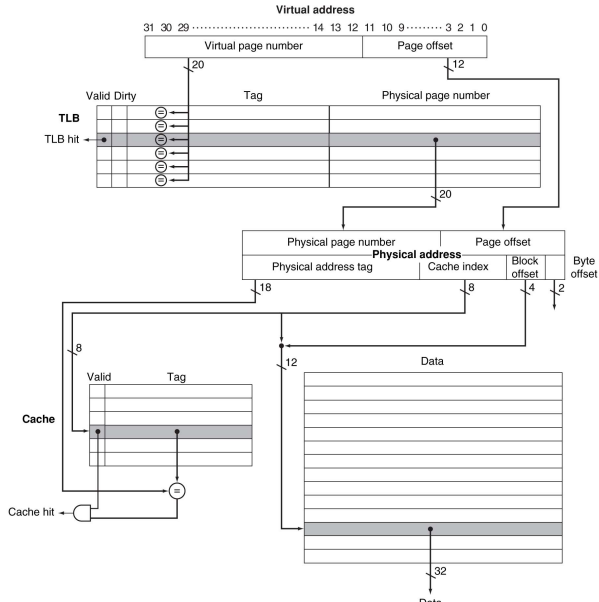
## Metrics – typical

- Size: typical 16-512 entries
- Block size: 1-2 page table entries (4-8 bytes each)
- Hit time: 0.5-1 cycle
- Miss penalty: 10-100 cycles (page hit)
- Miss rate: 0.01%-1%
- Levels: 1-2
- Split: data and instruction split on first level or both

## Associativity

- TLB's are small
- Need to be fast
- Fully associative OK for small caches
  - But (proper) LRU too expensive
- Rarely low associativity
  - Miss expensive  $\Rightarrow$  aim for low miss rate

# TLB and cache



# TLB examples

Characteristic	Intel Nehalem	AMD Opteron X4 (Barcelona)
Virtual address	48 bits	48 bits
Physical address	44 bits	48 bits
Page size	4 KB, 2/4 MB	4 KB, 2/4 MB
TLB organization	<p>1 TLB for instructions and 1 TLB for data per core</p> <p>Both L1 TLBs are four-way set associative, LRU replacement</p> <p>The L2 TLB is four-way set associative, LRU replacement</p> <p>L1 I-TLB has 128 entries for small pages, 7 per thread for large pages</p> <p>L1 D-TLB has 64 entries for small pages, 32 for large pages</p> <p>The L2 TLB has 512 entries</p> <p>TLB misses handled in hardware</p>	<p>1 L1 TLB for instructions and 1 L1 TLB for data per core</p> <p>Both L1 TLBs fully associative, LRU replacement</p> <p>1 L2 TLB for instructions and 1 L2 TLB for data per core</p> <p>Both L2 TLBs are four-way set associative, round-robin</p> <p>Both L1 TLBs have 48 entries</p> <p>Both L2 TLBs have 512 entries</p> <p>TLB misses handled in hardware</p>

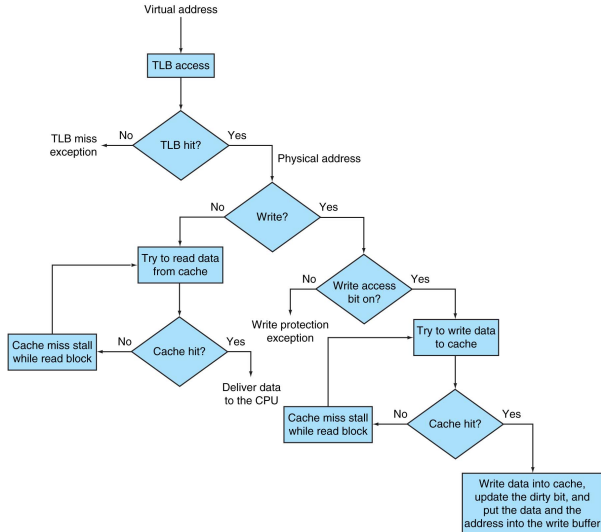
# TLB examples 2015-2017

	Intel Skylake (architecture) (2015/16)	AMD Zen (architecture) (2017)
Virtual address	48 bits (?)	?
Pyhsical address	- (depends on model)	- (depends on model)
Page size	4 KB, 2/4 MB, 1 GB	4 KB, 2/4 MB, 1 GB
TLB organization		
TLB L1 instruction	128 (4 KB pages, 8-way), 8 per thread (2/4 MB)	L0:8(all sizes); 64 (all sizes)
TLB L1 data	64 (4 KB pages), 32 (2/4 MB), 4 (1 GB); 4-way	64 (all sizes)
TLB L2	1536 (4 KB, 2/4 MB; 12-way), 16 (1G; 4-way)	I:512 (no 1G); D: 1536 (no 1G): 6-way

## TLB miss

- Not the same as page fault!
- Page hit
  - Can be handled by hardware or software
- Page fault
  - Handled by software
- Software is invoked via [exceptions](#)

# Memory access procedure



# Performance impact

- Working set  $>$  TLB covered address space
  - High TLB miss rate

# Performance impact

- Working set  $>$  TLB covered address space
  - High TLB miss rate

## Possible solution

- Larger pages
  - E.g. 2/4 MByte pages on x86



# Performance impact

- Working set  $>$  TLB covered address space
  - High TLB miss rate

## Possible solution

- Larger pages
  - E.g. 2/4 MByte pages on x86
- But, access of one byte pulls in 4 MB !!
  - Very bad for random access of large working set

# Performance impact

- Working set  $>$  TLB covered address space
  - High TLB miss rate

## Possible solution

- Larger pages
  - E.g. 2/4 MByte pages on x86
- But, access of one byte pulls in 4 MB !!
  - Very bad for random access of large working set

## Example:

- Radix Sort has more TLB misses than Quicksort
  - Slower despite the theoretical advantage

# TLB performance example

## Exercise

Determine TLB miss rate for code that goes one by one through a large array of double (8 bytes).

- Page size: 4KByte

## Exercise

Determine TLB miss rate for code that goes one by one through a large array of double (8 bytes).

- Page size: 4KByte

## Solution

- 4KByte page  $\Rightarrow$  512 array elements per page
- Access to first element: TLB miss
- Next 511 elements: TLB hit
- TLB miss rate =  $\frac{1}{512} = 0.195\%$

1 TLB

2 OS and Protection

# Role of kernel/OS

## Only kernel/OS

- Can write page tables
  - Sets the mapping between virtual and physical memory
- Sets page table register
  - Changed during context switch
- Can only be done in **supervisor/kernel mode**
  - Enter with **system call** or **exception**

# Role of kernel/OS

## Only kernel/OS

- Can write page tables
  - Sets the mapping between virtual and physical memory
- Sets page table register
  - Changed during context switch
- Can only be done in **supervisor/kernel mode**
  - Enter with **system call** or **exception**
- Context switch
  - Change of execution program (process)
  - TLBs and caches are cleared
    - *Can* make switching expensive
- Part of memory space must always be in physical memory
  - Exception handler
  - OS page table

# Page fault handler

- Part of kernel/OS
- Entered through exception



# Page fault handler

- Part of kernel/OS
  - Entered through exception
- 1 Find usable physical page
    - LRU algorithm
  - 2 Write it back to disc ([swap file/partition](#)) (if dirty bit set)

# Page fault handler

- Part of kernel/OS
  - Entered through exception
- 1 Find usable physical page
    - LRU algorithm
  - 2 Write it back to disc ([swap file/partition](#)) (if dirty bit set)
  - 3 Read requested page from disc

# Page fault handler

- Part of kernel/OS
  - Entered through exception
- 1 Find usable physical page
    - LRU algorithm
  - 2 Write it back to disc ([swap file/partition](#)) (if dirty bit set)
  - 3 Read requested page from disc
  - 4 Adjust page table entries

- Part of kernel/OS
  - Entered through exception
- 1 Find usable physical page
    - LRU algorithm
  - 2 Write it back to disc ([swap file/partition](#)) (if dirty bit set)
  - 3 Read requested page from disc
  - 4 Adjust page table entries
  - 5 Memory access retried
    - Cache miss ...

Virtual memory is used for protection

- Each program (process) has own page table
  - ⇒ Cannot read memory of other process
    - Simply not in their page table
    - Not allowed to change page table
- Sharing of memory must be explicitly enabled by OS
  - Different virtual addresses pointing to the same physical address

Virtual memory is used for protection

- Each program (process) has own page table
  - ⇒ Cannot read memory of other process
    - Simply not in their page table
    - Not allowed to change page table
- Sharing of memory must be explicitly enabled by OS
  - Different virtual addresses pointing to the same physical address
- **Protection bit** (read only)
  - Exception when written to
  - Ideal to share libraries
- **Execution bit**
  - Do not allow execution of data areas ⇒ **security**
- Exception handlers are very flexible and versatile
  - Can be used to implement transparent memory sharing over network!