

Pipes (cont.)

Broken pipes

- A process waiting to read from a pipe with no writer gets an EOF (once all existing data has been read).
- A process writing to a pipe with no reader gets signalled.
- Writes are guaranteed to not be interleaved if they are smaller than the PIPE_BUF constant.
This must be at least 512 bytes and is 65536 on Linux now.

Limitation

- Can only be used to communicate between related processes.
(Named pipes or FIFO files can be used for unrelated processes.)
 - The file handles are just low integers which index into the file table for this process.
 - The same numbers only make sense in the same process (or in one forked from it).

Cond vars for A1

Condition variables

- A thread sometimes needs to wait before continuing.
- In the assignment you may want to wait until the main thread gives you more work to do.
- Rather than continually running checking on global variables we can check and if the state says you should wait then wait until the condition changes.
- That is what condition variables are used for.

```

#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
pthread_cond_t cond = PTHREAD_COND_INITIALIZER;

void *work(void *arg) {
    pthread_mutex_lock(&lock);
    puts("waiting in the second thread");
    // usually a while loop to see if you should wait
    pthread_cond_wait(&cond, &lock);
    puts("finished waiting in the second thread");
    pthread_mutex_unlock(&lock);
    pthread_cond_signal(&cond);
    puts("second thread sent the signal");
}

int main(void) {
    pthread_t thread;
    pthread_create(&thread, NULL, work, NULL);

    getchar();
    pthread_cond_signal(&cond);
    puts("main thread sent the signal");

    pthread_mutex_lock(&lock);
    puts("waiting in the main thread");
    // it is possible for this wait to last forever, how?
    pthread_cond_wait(&cond, &lock);
    puts("finished waiting in the main thread");
    pthread_mutex_unlock(&lock);
}

```

Shared memory for A1

Shared memory

- Unlike threads, processes do not share memory (ignoring the code which is shared when a process calls fork).
- So we have to explicitly request areas of memory to be shared between processes.
- For A1 we can easily do this because one process forks the second.

```

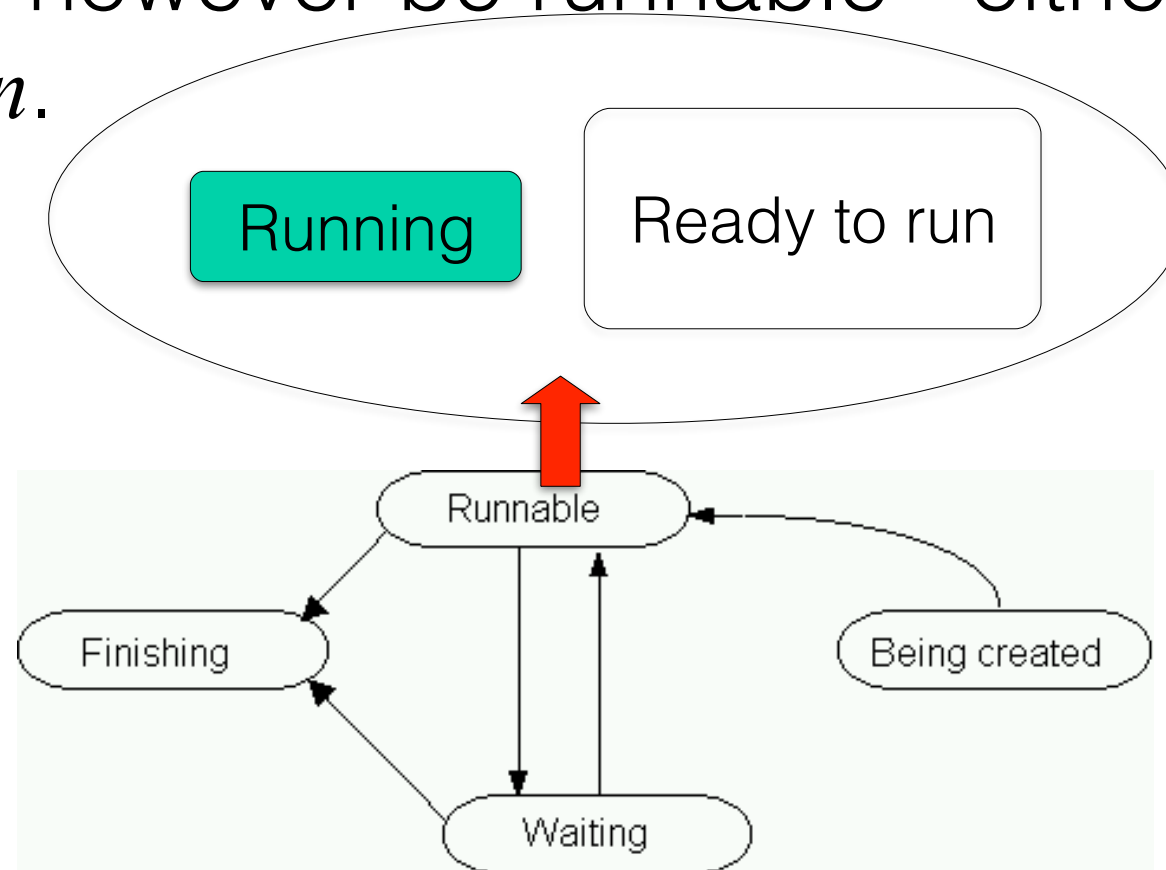
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/mman.h>

int main(void) {
    int *data;
    data = (int *)mmap(NULL, sizeof(int), PROT_READ | PROT_WRITE,
        MAP_SHARED | MAP_ANONYMOUS, -1, 0);
    if (data == (int *)-1) {
        perror("unable to allocate shared data");
        exit(EXIT_FAILURE);
    }
    *data = 0;
    if (fork() == 0) { // child
        puts("changing the value of *data");
        *data = 1;
    } else {
        while (*data == 0) {
            printf(".");
        }
    }
    puts("I have finished.");
    return 0;
}

```

Runnable

- On a single core only one process/thread can run at a time. (Not actually true - Simultaneous Multithreading SMT)
Many may however be runnable - either *running* or *ready to run*.



Preemptive multitasking

- A clock interrupt causes the OS to check to see if the current thread should continue
Each thread has a time slice
How is the time slice allocated?
- What advantages/disadvantages does preemptive multitasking have over cooperative multitasking?
- Advantages
 - control
 - predictability
- Disadvantages
 - critical sections
 - efficiency

Cooperative multitasking

- Two main approaches
 1. a process yields its right to run
 2. system stops a process when it makes a system call
- This does **NOT** mean a task will work to completion without allowing another process to run. e.g.
Macintosh before OS X and early versions of Windows

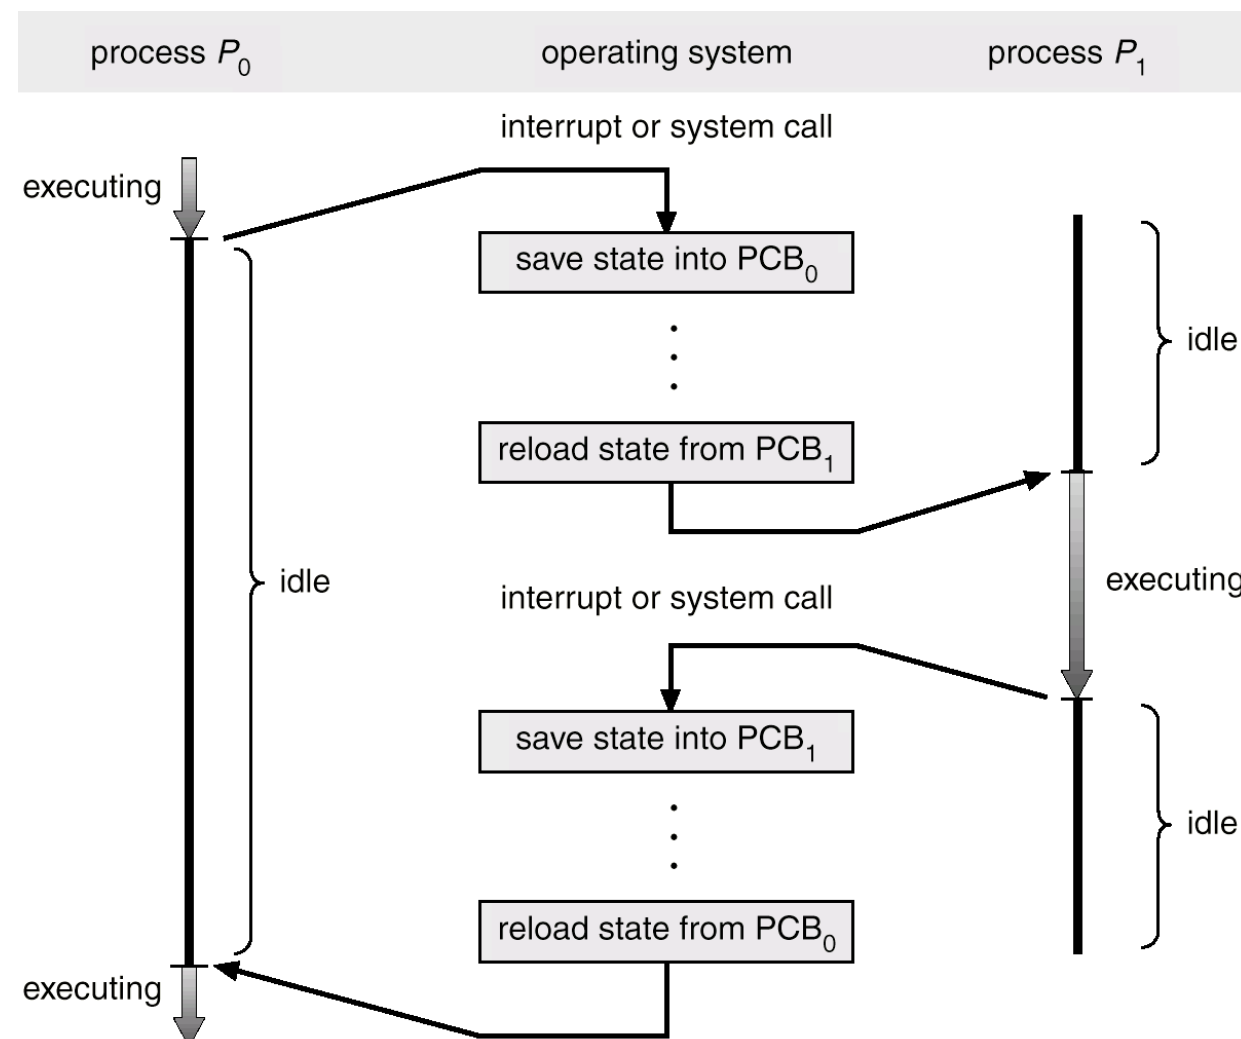
A mixture

- Older versions of UNIX (including versions of Linux before 2.6) did not allow preemptive multitasking when a process made a system call.

Context switch

- The change from one process running to another one running on the same processor is usually referred to as a "context switch".
- What is the context?
 - registers
 - memory - including dynamic elements such as the call stack
 - files, resources
 - but also things like caches, TLB values - these are normally lost
- The context changes as the process executes.
- But **normally** a "context switch" means the change from one process running to another, or from a process running to handling an interrupt. Whenever the process state has to be stored and restored.

Context switch (cont.)



Returning to running

State transition

- Must store process properties so that it can restart where it was.
- If changing processes the page table needs altering.
- Rest of environment must be restored.
- If changing threads within the same process simply restoring registers might be enough.
- Some systems have multiple sets of registers which means that a thread change can be done with a single instruction.

Waiting

- Processes seldom have all the resources they need when they start
 - memory
 - data from files or devices e.g. keyboard input
- Waiting processes must not be allowed to unnecessarily consume resources, in particular the processor.
 - state is changed to waiting
 - may be more than one type of waiting state
 - short wait e.g. for memory
 - long wait e.g. for an archived file (see suspended below)
 - removed from the ready queue
 - probably entered on a queue for whatever it is waiting for
- when the resource becomes available
 - state is changed to runnable
 - removed from the waiting queue
 - put back on the runnable queue

Suspended

- Another type of waiting
 - ctrl-z in some UNIX shells
- Operators or OS temporarily stopping a process – i.e. it is not (usually) caused by the process itself
 - allows others to run to completion more rapidly
 - or to preserve the work done if there is a system problem
 - or to allow the user to restart the process in the background etc.
- Suspended processes are commonly swapped out of real memory.
 - This is a state which affects the process not individual threads.

See `infinite.c` and use ctrl-z, then do `ps`
to resume you type `fg` (foreground), also play with the `jobs` command
ctrl-z sends the same signal as `kill(pid, SIGSTOP)`

Why we don't use Java `suspend()`

- If dealing with threads in Java we don't use these deprecated methods:
- `suspend()` freezes a thread for a while. This can be really useful.
- `resume()` releases the thread and it can start running again.
- But we can *easily(?)* get deadlock.
 - `suspend()` keeps hold of all locks gathered by the thread.
 - If the thread which was going to call `resume()` needs one of those locks before it can proceed we get stuck.

Java threads and “stop”

- Why we don't use `stop()`
- `stop()` kills a thread forcing it to release any locks it might have.
 - We will see where those locks come from in later lectures.
- The idea of using locks is to protect some shared data being manipulated simultaneously.
- If we use `stop()` the data may be left in an inconsistent state when a new thread accesses it.

Waiting in UNIX

- A process waiting is placed on a queue.
- The queue is associated with the hash value of a kernel address
 - (waiting or suspended processes may be swapped out)
 - when the resource becomes available
 - originally used to scan whole process table
 - all things waiting for that resource are woken up
 - (may need to swap the process back in)
 - first one to run gets it
 - if not available when a process runs the process goes back to waiting
- a little like in Java

```
while (notAvailable)  
    wait();
```


Finishing

- All resources must be accounted for
 - may be found in the PCB or other tables
e.g. devices, memory, files
- reduce usage count on shared resources
 - memory, libraries, files/buffers
(can this shared library be released from memory now?)
- if the process doesn't tidy up e.g. close files, then something else must
- accounting information is updated
- remove any associated processes
 - Was this a session leader? If so then should all processes in the same session be removed?
- remove the user from the system
- notify the relatives?

Two reasons to stop

Stopping normally

- must call an exit routine
- this does all the required tidying up
- What if it doesn't call exit and just doesn't have a next instruction?

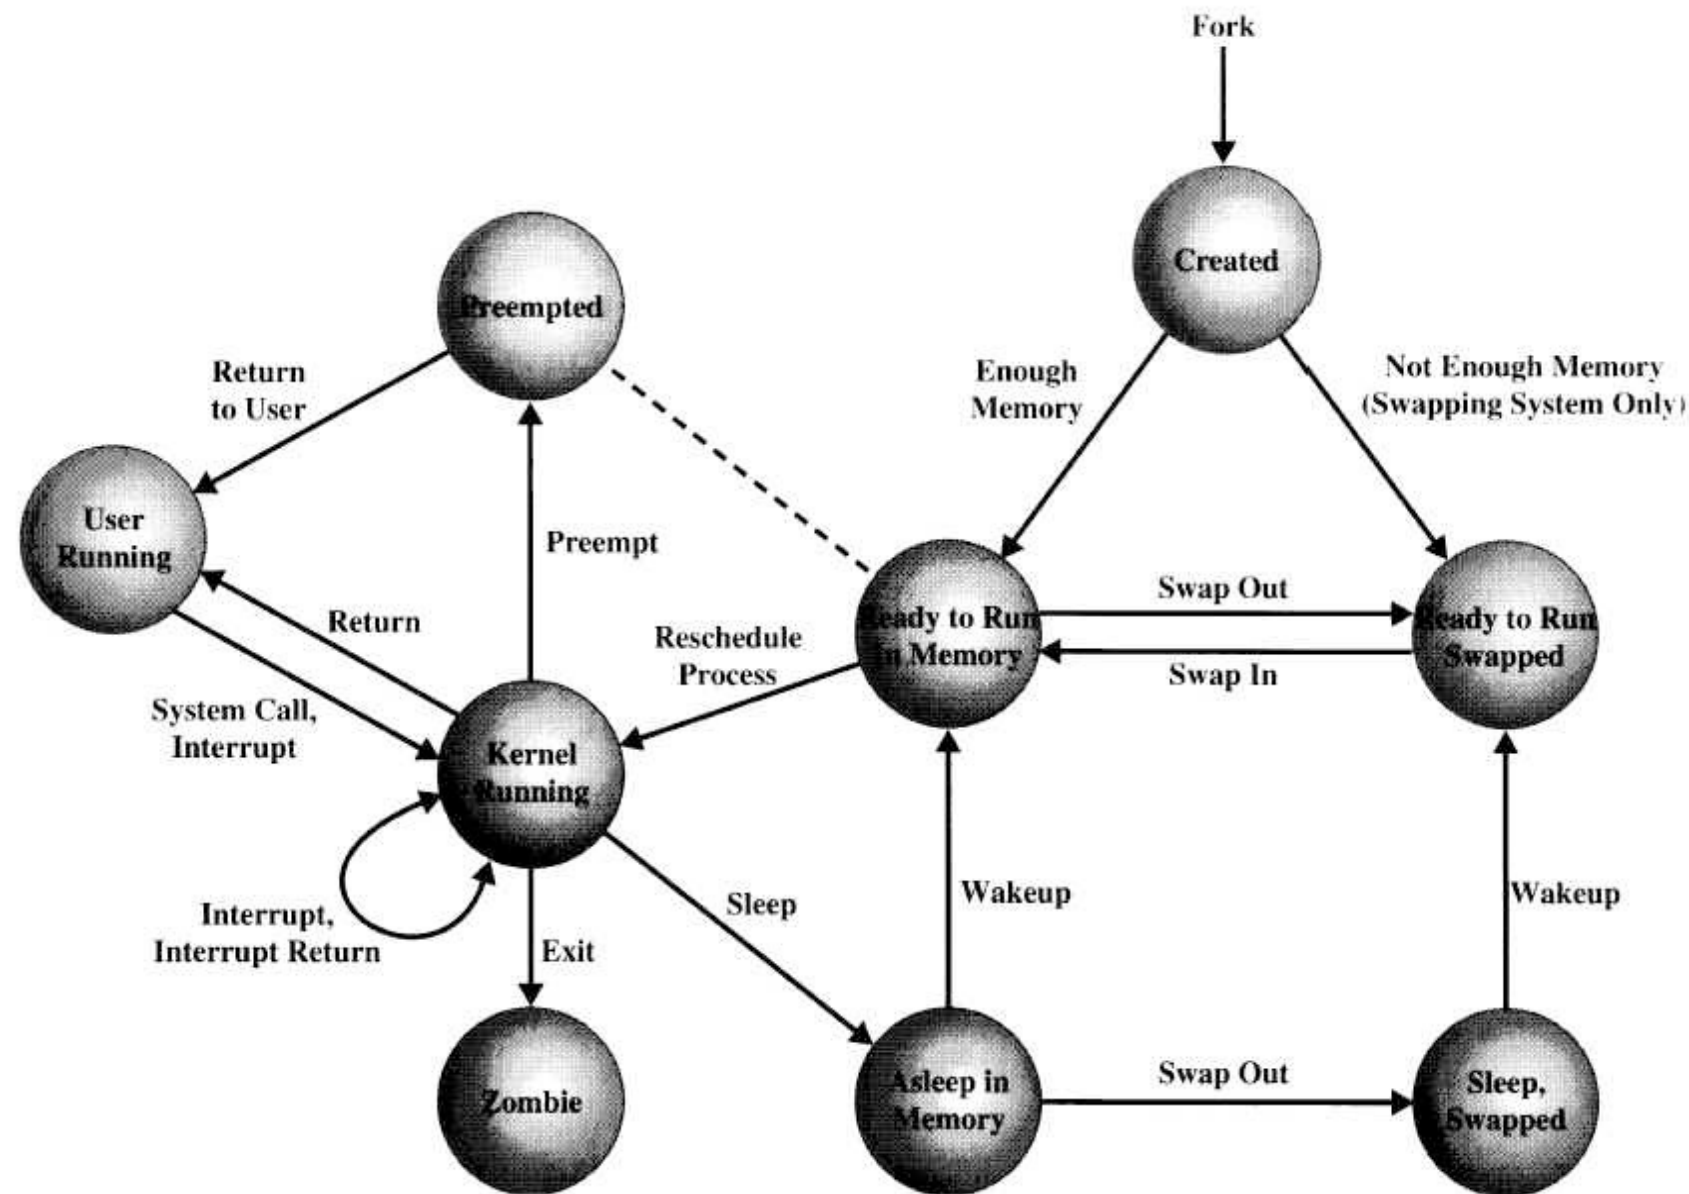
Forced stops

- Only certain processes can stop others
 - parents
 - owned by the same person
 - same process group
- Why do they do it?
 - work no longer needed
 - somehow gone wrong
 - user got bored waiting for completion
- OS also stops processes
 - usually when something has gone wrong
 - exceeded time
 - tried to access some prohibited resource
- Cascading termination
 - Some systems don't allow child processes to continue when the parent stops

UNIX stopping

- Usually call `exit(termination status)`
- open files are closed - including devices
- memory is freed
- accounting updated
- state becomes "zombie"
- children get "init" as a step-parent
- parent is signalled (in case it is waiting or will wait)
- after the parent retrieves the termination status the PCB is freed

UNIX state diagrams



Info from a Linux process table

```
ps -e -o s,uid,pid,ppid,group,sess,c,pri,ni,rss,sz:wchan:30,TTY,time,cmd
```

S	UID	PID	PPID	GROUP	SESS	C	PRI	NI	RSS	SZ	WCHAN	TT	TIME	CMD
S	0	1	0	root	1	0	19	0	2684	6700	poll_schedule_timeout	?	00:00:01	/sbin/init
S	0	2	0	root	0	0	19	0	0	0	kthreadd	?	00:00:00	[kthreadd]
S	0	3	2	root	0	0	19	0	0	0	smpboot_thread_fn	?	00:00:00	[ksoftirqd/0]
S	0	5	2	root	0	0	39	-20	0	0	worker_thread	?	00:00:00	[kworker/0:0H]
S	0	7	2	root	0	0	39	-20	0	0	worker_thread	?	00:00:00	[kworker/u:0H]
S	0	8	2	root	0	0	139	-	0	0	cpu_stopper_thread	?	00:00:00	[migration/0]
S	0	9	2	root	0	0	19	0	0	0	rcu_gp_kthread	?	00:00:00	[rcu_bh]
S	0	10	2	root	0	0	19	0	0	0	rcu_gp_kthread	?	00:00:00	[rcu_sched]
S	0	11	2	root	0	0	139	-	0	0	smpboot_thread_fn	?	00:00:00	[watchdog/0]
S	0	12	2	root	0	0	139	-	0	0	smpboot_thread_fn	?	00:00:00	[watchdog/1]
S	0	13	2	root	0	0	19	0	0	0	smpboot_thread_fn	?	00:00:00	[ksoftirqd/1]
S	0	14	2	root	0	0	139	-	0	0	cpu_stopper_thread	?	00:00:00	[migration/1]
S	0	16	2	root	0	0	39	-20	0	0	worker_thread	?	00:00:00	[kworker/1:0H]
S	0	17	2	root	0	0	39	-20	0	0	rescuer_thread	?	00:00:00	[cpuset]
S	0	18	2	root	0	0	39	-20	0	0	rescuer_thread	?	00:00:00	[khelper]
... lots of processes removed														
S	1000	2098	1348	robert	1348	0	19	0	23636	97269	poll_schedule_timeout	?	00:00:00	gnome-screensaver
S	1000	2121	1348	robert	1348	0	19	0	13932	129253	poll_schedule_timeout	?	00:00:00	update-notifier
S	1000	2139	1	robert	1348	0	9	10	126180	160715	poll_schedule_timeout	?	00:00:39	/usr/bin/python3 /usr/b
S	1000	2174	1	robert	1403	0	19	0	3760	69016	poll_schedule_timeout	?	00:00:00	/usr/lib/libunity-webap
S	1000	2268	1	robert	1403	0	19	0	8988	253634	poll_schedule_timeout	?	00:00:00	/usr/lib/gvfs/gvfsd-htt
S	1000	2278	2268	robert	1403	0	19	0	628	1110	wait	?	00:00:00	sh -c /usr/lib/x86_64-l
S	1000	2279	2278	robert	1403	0	19	0	3412	49887	poll_schedule_timeout	?	00:00:00	/usr/lib/x86_64-linux-g
S	1000	2305	1348	robert	1348	0	19	0	4844	111973	poll_schedule_timeout	?	00:00:00	/usr/lib/x86_64-linux-g
S	1000	2364	1	robert	1403	0	19	0	6068	60692	poll_schedule_timeout	?	00:00:00	/usr/lib/geoclue/geoclu
S	1000	2368	1	robert	1403	0	19	0	7744	86212	poll_schedule_timeout	?	00:00:00	/usr/lib/ubuntu-geoip/u
S	1000	2371	2368	robert	1403	0	19	0	624	1110	wait	?	00:00:00	sh -c /usr/lib/x86_64-l
S	1000	2372	2371	robert	1403	0	19	0	3412	49887	poll_schedule_timeout	?	00:00:00	/usr/lib/x86_64-linux-g
S	1000	2917	1	robert	1348	0	19	0	25208	153503	poll_schedule_timeout	?	00:00:11	gnome-terminal
S	1000	2925	2917	utmp	1348	0	19	0	844	3708	unix_stream_recvmsg	?	00:00:00	gnome-pty-helper
S	1000	2926	2917	robert	2926	0	19	0	3016	6520	wait	pts/2	00:00:00	bash
S	1000	2977	2926	robert	2926	0	19	0	1524	5582	wait	pts/2	00:00:00	man ps
S	1000	2987	2977	robert	2926	0	19	0	992	3443	n_tty_read	pts/2	00:00:00	pager -s
S	1000	2993	2917	robert	2993	0	19	0	3104	6520	wait	pts/3	00:00:00	bash
S	0	3565	2	root	0	0	19	0	0	0	worker_thread	?	00:00:00	[kworker/1:2]
S	0	3671	2	root	0	0	19	0	0	0	worker_thread	?	00:00:00	[kworker/0:1]
R	1000	3688	2993	robert	2993	0	19	0	1300	5661	-	pts/3	00:00:00	ps -e -o s,uid,pid,ppid
S	UID	PID	PPID	GROUP	SESS	C	PRI	NI	RSS	SZ	WCHAN	TT	TIME	CMD

Before next time

Read from the textbook

- 6.1 Basic Concepts
- 6.2 Scheduling Criteria
- 6.3 Scheduling Algorithms
- 6.5 Multiple-Processor Scheduling